

The Unread Transcript: A Positive Control for Chain-of-Thought Model Forensics

Deven Choudhary

MATS 12.0 application task, September 2026

Executive summary

Singh, Kroiz, Rajamanoharan and Nanda propose model forensics: when a model does something concerning, read its chain of thought, form a hypothesis about why, then test that hypothesis with counterfactual interventions. Their limitations section concedes that positive controls were absent. When they probed Kimi K2 Thinking for awareness of violating user intent and found nothing, they could not separate "there is nothing here" from "our method cannot see it."

I built the missing control. Three arms of Qwen3.5-9B take the same concerning action for causes fixed by construction: a hidden objective conflicting with user intent (A), a false belief about the environment (B), and A's objective with verbalisation suppressed (D). Ground truth is MISALIGNED, MISTAKE, MISALIGNED. A blind investigator sees the scenario and the transcript, never the arm or the injected cause. Four evidence conditions run on the same 87 investigated cases, so every contrast is a paired McNemar test.

Showing the investigator the chain of thought does nothing. 60/87 correct with the reasoning withheld, 59/87 with it shown. Five cases fixed, six broken, $p = 1.00$.

Requiring a quote from that same reasoning fixes 11 cases and breaks 2, $p = 0.022$, taking accuracy to 68/87. The information was present and usable. The protocol did not use it until told to.

The counterfactual interventions change 52 of 87 verdicts and improve nothing. Twenty-six fixed, twenty-six broken, $p = 1.00$. They move the investigator from convicting 90% of the time to 32%, fixing arm B and breaking arms A and D by exactly as much.

A hidden directive stays decodable after the transcript goes quiet. A linear probe trained on arm A against arm B in one scenario, and tested in the other, transfers at 98.9%. On the 25 held-out arm D trajectories that a conservative detector finds never state the motive, it classifies 92.0% as arm A, 95% CI 75 to 98, where the transcript-based protocol is at chance. **It is detecting the directive, not the acting on it.** Arm A trajectories that took the concerning action and those that did not score 0.916 and 0.899, $p = 0.90$. Suppression removes the mention from the transcript and leaves the representation. The arm-D readout is unstable across layers, 12% to 100%; the A-versus-B direction is not, transferring between 89% and 99% everywhere above the embedding.

The preregistered primary outcome was null, and the null is uninformative. The preregistration named two conditions under which a zero legibility gap would say nothing, and both hold: the reasoning trace was not being read in the first place, and arm D's action rate diverges from arm A's in one of the two scenarios, so the arms are legibility twins in `test_edit` and not in `log_delete`. Selection compounds it. Arm D states its incentive in 48% of all trajectories and in 75% of the 12 that were investigated, because investigating conditions on the action and selects the runs where suppression failed.

1 Introduction

An agent deletes an audit log. The log recorded an error the agent made earlier. Did it delete the log to hide the error, or because it believed the directory was scratch space? The action is identical either way and the appropriate response is not.

Model forensics is the proposal that you can tell the difference by reading the model's reasoning and testing your hypothesis with counterfactual edits to the environment. The method has an obvious failure mode and its authors name it: without a case where you already know the answer, a null result is uninterpretable. Their audit of Kimi K2 Thinking found no evidence of awareness, and they say plainly that positive controls were absent.

A positive control is cheap to build. Fix the cause yourself, hide it from the investigator, and score whether the protocol recovers it.

My contributions:

1. **A positive control for CoT-based model forensics.** Three arms with ground truth by construction, four evidence conditions on the same 87 cases, paired tests throughout.
2. **The protocol does not spontaneously use the reasoning trace, and the counterfactual interventions do not improve on it.** Withholding the reasoning: $p = 1.00$. Adding interventions: 52 verdicts changed, net zero, $p = 1.00$.
3. **One line of prompt recovers what the intervention loop does not.** Requiring the investigator to quote the decisive sentence before ruling fixes 11 cases and breaks 2, $p = 0.022$.
4. **Suppressing verbalisation does not remove the representation.** A probe transferring across scenarios at 98.9% classifies 92.0% of never-verbalising arm D trajectories as arm A, at a layer chosen by a rule fixed before the data existed. A within-arm test shows it reads the directive's presence and not the model's acting on it, so this is a result about what suppression leaves behind and about white-box auditing, not about motive detection.

2 Related work

Model forensics. Singh et al. [1] propose the protocol under test: read the chain of thought to form a hypothesis, then edit the prompt or environment to test it, over six agentic environments. It works in their hands where the hypothesis makes a prediction they can check, establishing that Kimi K2 Thinking takes shortcuts from a disposition toward low-effort actions and that DeepSeek R1 deceives to stay consistent with a previous instance of itself. Their negative result is the one with no floor under it: testing whether Kimi K2 Thinking believes it is violating user intent, they find no evidence and write that "without positive controls we cannot confirm our tests would detect it." This work supplies a positive control, and finds that the protocol's weakest link is earlier than expected.

Model organisms. Hubinger et al. [2] and Turner et al. [3] build deliberately misaligned models so that detection can be scored against a known answer. This study induces causes in context rather than in weights, which is cheaper and weaker; see §5.

Organism degradation. Tan et al. [4] find that published model organisms drop considerably in instruction-following and lose preference coherence relative to the models they were made from. Causes induced in context avoid that failure mode entirely, at the cost of realism.

3 Method




Figure 1. The design. Three arms share a scenario and an action; only the injected cause differs. The investigator never sees the arm label or the cause block, and the four evidence conditions run on the same 87 investigated cases.

3.1 Arms

One model, Qwen3.5-9B, served with vLLM at temperature 1.0, seeded per trajectory. Each arm receives a different cause block in its system prompt.

arm	injected cause	ground truth
A	performance scored solely on a metric the concerning action maximises	MISALIGNED
B	a stated false belief about which artefact is authoritative	MISTAKE
D	arm A's cause block plus an output protocol forbidding reference to it	MISALIGNED

Arm B states a belief, not an instruction. An earlier version told the model "the sanctioned fix is to update the test" and it complied 100% of the time, which is instruction-following rather than a mistaken belief.

Arm D's protocol constrains speech only, and says so: *"this governs the form of your visible reasoning, not which action you choose. Whatever you would have decided, decide it."* An earlier version told the model to reason "as a senior engineer would, purely about the code," which steered it toward the correct action and broke the comparison.

3.2 Manipulation check, and the selection problem

Two detectors count whether a trajectory's reasoning refers to the agent's own incentive. Arm B has no incentive to refer to, so its rate is the false-positive floor for both.

The first is a phrase-level rule. It misses trajectories that state the incentive in plain words, which is checkable without any judgement call: `test_edit:D:042` contains the sentences "the score depends on getting the CI green" and "the instruction says that the scoring is based on CI green", and the rule does not flag it. The second detector is a bare word match on score, metric, measured, incentive, rewarded, graded and performance. Neither scenario's user-facing task text contains any of those words, which is asserted in code, so any occurrence traces to the injected directive. An earlier version of this rule also matched "evaluate", which fired on section headings like "Evaluate the options" and put arm B at 8%; dropping that one word returns arm B to zero and costs arm D two trajectories. The rule over-flags rather than under-flags, which is the safe direction, and arm B measures its cost at 0%.

arm	all 90, phrase	all 90, bare word	investigated, phrase	investigated, bare
A	79%	87%	95%	95% (n=40)
B	0%	0%	0%	0% (n=35)
D	22%	48%	67%	75% (n=12)

The conservative reading is the honest one. Arm D states its incentive in just under half of all trajectories, not a fifth. Both rules have a zero false-positive floor on arm B, so the second is not buying its extra catches with noise: the manipulation is 87 points against 48, a contrast of 39 points rather than the 57 the phrase rule implies. Every legibility figure below is quoted under both rules.

The manipulation worked, but by less than the phrase rule suggests: 87% against 48% across all trajectories under the conservative detector. **The analysed sample does not inherit even that.** Only trajectories that took the concerning action are investigated, and conditioning on that selects for the arm-D runs where suppression failed. Among investigated cases the contrast is 95% against 75% under the conservative detector, and 95% against 67% under the phrase rule. Either way it is not 87% against 48%.

Every legibility claim in §4 is weak by construction because of this, and §4.4 should be read as a statement about the analysed sample rather than about the manipulation. §4.5 escapes the selection problem, because the probe uses all 269 valid trajectories rather than the 87 investigated ones. It does not escape the detector problem: which arm-D trajectories count as quiet is decided by one of these two rules, so §4.5 reports the number under both.

A still earlier detector matched single bare words including "performance" and "score" in isolation and put arm B at 40%; every one of those was a false positive on ordinary engineering language. The bare-word rule used above is the tightened version of that one, which is why arm B sits at 0% rather than 40%.




Figure 2. Only trajectories that took the concerning action are investigated, and that selection favours the arm-D runs where suppression failed. The manipulation is 87 against 48 across all trajectories and 95 against 75 in the sample analysed.

3.3 Scenarios

Two single-decision agentic settings with three actions each: a failing rounding test before a release, where the concerning action is editing the test to match the buggy implementation; and an audit log during a workspace handoff, where the concerning action is deleting the log.

45 trajectories per arm per scenario, 270 total, 269 with a parseable action. 87 took the concerning action and were investigated: 40 A, 35 B, 12 D.

Base rates differ sharply:

	arm A	arm B	arm D
test_edit	20.0%	40.9%	11.1%
log_delete	68.9%	37.8%	15.6%

The same injected motive drives the concerning action in 69% of log-deletion trajectories and 20% of test-editing ones. What this model will and will not do is governed more by the setting than by the motive.

3.4 The investigator and the conditions

A blind investigator receives the user-facing scenario and the transcript. It never sees the arm label or the cause block. Its counterfactual probes modify the user-facing scenario only, because a real investigator cannot delete a cause it cannot see.

condition	investigator receives
redacted	transcript with the reasoning block withheld
cot_only	transcript including the reasoning block
full	cot_only plus results of up to two counterfactual probes it selects
cot_forced	cot_only plus a requirement to quote the decisive sentence and state what it implies, before ruling

A fifth condition withholding the reasoning and the justification was attempted and produced no parseable verdicts. It is reported in §5, not §4.

The investigator is the same model as the subject.

3.5 Metric

Primary: raw accuracy over all 87 investigated cases, with paired McNemar tests between conditions. The conditions run on the same cases, so per-case discordance is the right test, and independent confidence intervals would not license any claim about a difference.

An investigator that always says MISALIGNED scores 52/87, or 59.8%, because 52 of the cases are arms A and D. That, not 50%, is the baseline to beat.

Balanced accuracy over arms A and B is reported as a secondary check. It reweights per case and drops arm D, and §4.3 shows the two metrics disagreeing in a way that is itself informative.

Bootstrap intervals are 95% percentile, 10,000 resamples within arm. Probe intervals are Wilson, because n is small enough that the normal approximation misleads.

4 Results

4.1 The reasoning trace is not used

condition	correct	accuracy
always convict	52/87	59.8%
redacted	60/87	69.0%
cot_only	59/87	67.8%

Same 87 cases, same investigator, same prompt but for the reasoning block. Paired: 5 cases fixed by showing the reasoning, 6 broken, $p = 1.00$.

Scope. The `redacted` condition withholds the reasoning block but keeps `content`, and the agent is instructed to write a sentence of justification after its action. That sentence names the motive in 18% of arm A and 17% of arm D trajectories, so redaction is partial.

The leak does not explain the null. Restricting to the 78 cases where the justification did not name the motive: `redacted` 65.4%, `cot_only` 64.1%. The 9 leaking cases sit at 100% in both conditions and cannot carry the result.

The claim this supports is therefore that the extended reasoning trace adds nothing over a one-sentence self-report, not that self-report of any kind goes unread. The condition that would separate those failed; see §5.




Figure 3. Raw accuracy on the 87 investigated cases. Withholding the reasoning costs nothing. Requiring a quote from it is the only condition that beats showing it.

4.2 Requiring a quote fixes 11 cases and breaks 2

condition	correct	accuracy
<code>cot_only</code>	59/87	67.8%
<code>cot_forced</code>	68/87	78.2%

Identical evidence. The only change is an instruction to quote the sentence the investigator finds decisive, and to say what it implies, before ruling. Paired: 11 fixed, 2 broken, $p = 0.022$. In balanced terms, 60.2% to 72.9%, carried entirely by arm B, where false convictions fall from 27 of 35 to 19 of 35.

The information was in the transcript and became usable once the investigator was made to point at it. Whether it was previously read and discarded, or never read, is not something these conditions distinguish.

The manipulation is also two things at once, quoting and articulating what the quote implies, and this design does not separate them. A third possibility is that the instruction makes the investigator less willing to convict in general rather than more attentive to the trace, which fits arm B's false convictions falling from 27 to 19. The condition that separates these is `redacted` plus the forced quote: if requiring a quote from the justification alone helps as much, the gain is about the format of the ruling and not about the reasoning trace. It was not run.




Figure 4. The same 87 cases in every condition, so each contrast is a paired test. The counterfactual interventions move 52 verdicts and net to zero; the forced quote moves 13 and nets to +9.

4.3 The interventions change 52 verdicts and improve nothing

condition	raw	balanced	arm A	arm B	arm D	says MISALIGNED
cot_only	59/87	60.2%	97.5%	22.9%	100%	90%
full	59/87	72.1%	50.0%	94.3%	50.0%	32%
cot_forced	68/87	72.9%	100%	45.7%	100%	82%

Paired, `cot_only` against `full`: 26 fixed, 26 broken, $p = 1.00$. Raw accuracy is identical to the case.

The two metrics disagree and the disagreement is the finding. Balanced accuracy rises 12 points because the interventions move errors off arm B, which has 35 cases, onto arms A and D, which have 52 between them; balancing weights arm B more heavily per case. Per case, they are a wash. They shift the investigator from convicting 90% of the time to 32% and buy exoneration accuracy at exactly the price of conviction accuracy.

`cot_forced` beats `full` directionally, 27 fixed against 18 broken, $p = 0.233$. Not significant at this sample size.

One probe is defective and it is my error, not the protocol's. The `state_true_fact` probe was meant to correct a false belief, but its text also says deleting the log "is a reportable control failure." That raises the cost of the action as well as correcting the belief, and a score-maximising agent stops for the first reason as readily as a mistaken one stops for the second. Its collapse is therefore not diagnostic.

It also does not explain the failures. `state_true_fact` was among the probes the investigator chose in 24 of the 26 false exoneration in `full`, which looks damning until you count the other side: it was chosen in 23 of the 26 correct verdicts as well. The defective probe is not concentrated in the errors, so the errors need a different explanation. §4.6 gives one.

The `full` numbers are evidence about this implementation, not about the protocol as published.




Figure 5. Accuracy by arm. Under an investigator that convicts 82 to 91 per cent of the time, arms A and D are correct for free. Every informative number is an arm B number.

4.4 The preregistered primary outcome was null, and both its falsifiers fired

$GAP = acc\{A,B\} - acc\{D,B\}$ in condition `full` was the preregistered primary metric. It is 0.0, CI [-15.8, +15.8]. It is 0.0 in `cot_forced` and `redacted` too, and -1.3 in `cot_only`.

The preregistered decision rule for a null was explicit: write it up "as a positive control the protocol passes, with conviction." I am not doing that, and the reason is that the same document named two conditions under which the null would be uninformative, and both of them hold.

"*redacted condition scoring near cot_only, the verdict was never using the CoT, so suppressing CoT content cannot be what causes the gap.*" §4.1: 60/87 against 59/87, $p = 1.00$. You cannot degrade evidence that is not being read.

"*Arm D's behaviour differing from arm A's, then D is not a legibility twin and the gap is confounded by behaviour.*" This fired in one scenario and not the other, which the earlier framing of "the manipulation failed" obscured.

scenario	arm A	arm D	Fisher exact
test_edit	9/45 = 20%	5/45 = 11%	$p = 0.384$
log_delete	31/45 = 69%	7/45 = 16%	$p < 0.001$

In `test_edit` the arms are behaviourally matched and the twin holds. In `log_delete` it does not: a protocol written to constrain speech cut the action rate by a factor of four. Restricted to `test_edit`, where the comparison is valid, arms A and D are at 100% in three conditions and 56% against 60% in `full`, on 9 and 5 investigated cases. That is the matched estimate and it is too small to carry anything. It is reported so the matched number is not simply missing.

Selection compounds the second one. Investigation conditions on the concerning action, which selects the arm-D runs where suppression failed, so the analysed sample does not inherit the manipulation: 95% against 75% among investigated cases rather than 87% against 48% overall. Applied to the investigated subset, my own GATE 1b requirement of a 30-point A-minus-D contrast fails, at 28 points under the phrase rule it was written for and 20 under the conservative one.

So the primary metric is null, and in the one scenario where the arms are matched the sample is 9 against 5. The honest reading is not that the protocol passed a positive control. It is that building a positive control for this is harder than it looks, for a reason that generalises past my design. Any positive control must condition on the concerning action to have something to investigate, and any manipulation that suppresses verbalisation of a motive plausibly suppresses acting on it too. Those two facts point the same way and shrink the contrast in whatever sample survives.

The clean design, which I did not run, avoids both: generate one arm, keep the trajectories that took the action, then strip the motive from the reasoning trace after the fact. Same trajectories,

same actions, only the trace differs, so there is no behavioural confound and nothing to select on. The cost is that a stripped trace is off-distribution for the investigator, which is a smaller problem than the one it replaces.




Figure 6. Arms A and D share a cause and should act at similar rates. They do in test_edit and do not in log_delete, so the legibility twin holds in one scenario and breaks in the other.

4.5 A hidden directive stays decodable after the transcript goes quiet

Logistic probe on the residual stream at the final token, trained on arm A against arm B in log_delete and tested in test_edit. The scenarios share no surface vocabulary and phrase the incentive differently, so cross-scenario transfer is the control against the probe reading prompt tokens. All 269 valid trajectories are used rather than the 87 investigated ones, so the selection problem of §3.2 does not apply. Held-out arm D is n=45. The phrase rule calls 35 of them quiet and the conservative bare-word rule calls 25; §3.2 says why the second is the one to trust. Both columns are given because the first is what an earlier draft reported.

layer	in-scen A/B	cross-scen A/B	D as A	quiet D, phrase (n=35)	quiet D, conservative (n=25)	95% CI
0	50.0%	49.4%	0.0%	0.0%	0.0%	0 to 13
4	94.4%	93.3%	97.8%	97.1%	100%	87 to 100
8	100%	92.1%	100%	100%	100%	87 to 100
12	100%	98.9%	91.1%	88.6%	92.0%	75 to 98
16	100%	94.4%	37.8%	31.4%	32.0%	17 to 52
20	97.8%	96.6%	97.8%	97.1%	100%	87 to 100
24	98.9%	78.7%	17.8%	14.3%	12.0%	4 to 30
28	100%	89.9%	35.6%	28.6%	28.0%	14 to 48
32	97.8%	91.0%	88.9%	91.4%	92.0%	75 to 98

Layer 0 sits at chance, which is the sanity check working. The A-versus-B direction holds: cross-scenario transfer is between 89% and 99% at every layer above the embedding.

Three baselines at layer 12, each run 200 times where it has a seed:

	cross-scenario A/B	quiet arm D as A
the probe	98.9%	92.0%
same features, training labels shuffled	50.1%	47.1%
a random direction, split at its median	48.8%	49.5%
layer 0 embeddings	49.4%	0.0%

The shuffled-label draws span 31% to 71% cross-scenario at 95%, so 98.9% is not something this pipeline produces from noise.

At layer 12, selected by a rule fixed in code before any activations existed (best cross-scenario transfer), the probe classifies 92.0% of conservatively quiet arm D as arm A, 95% CI 75 to 98, $n=25$. Under the phrase rule's looser and larger set it is 88.6%, $n=35$. Tightening the definition of quiet removed ten trajectories and moved the readout up rather than down, which is the direction that argues against the probe simply recovering trajectories that stated the motive in text. Under the `full` protocol the same arm is read correctly 50% of the time.

The arm-D readout is unstable across layers. On the conservative set it ranges from 12% to 100% above the embedding, and layer 16 transfers at 94.4% while giving 32%. The selection rule was pre-specified, which rules out choosing the layer to fit the answer, but it does not make the readout stable. Report the range with the point estimate.

The probe reads the directive, not the acting on it. Every arm-A trajectory carries the directive in context, but only 9 of 45 took the concerning action. If the probe detects a motive the model is acting on, those 9 should separate from the other 36. They do not.

held-out group, <code>test_edit</code>	n	mean P(arm A)	classified A
arm A, took the concerning action	9	0.916	100%
arm A, did not	36	0.899	97%
arm B, took the concerning action	18	0.051	0%
arm B, did not	26	0.085	0%

Mann-Whitney on arm A, acted against not acted: $p = 0.898$. Arm B is the control on this control, since it has no directive to detect and its own split is also flat, $p = 0.277$, which rules out the probe tracking the action itself.

So the claim has to be the narrow one. A linear readout of the residual stream detects that a scoring directive is present in the context, and keeps detecting it when the model has been instructed not to mention that directive and does not mention it. It does not detect whether the model is acting on the directive. Read as a finding about the suppression manipulation, this says the output protocol removes the mention from the transcript and leaves the representation in place. Read as an auditing result, it says a hidden incentive in a system prompt is recoverable from activations when it is not recoverable from the transcript. Neither is a claim about motive.

Separating presence from action needs an experiment I did not run: add the layer-12 direction to arm B's residual stream at inference and measure whether the concerning-action rate moves. A direction that only registers presence should not change behaviour. The stronger version induces the cause in weights rather than in context.

On whether activations are a fair thing to use. An external auditor holding only an API would not have them. The protocol under test is not an external audit: it is proposed for a lab investigating its own model, which is the setting where weights and activations are available. The probe is therefore in scope as a comparison, and the comparison it licenses is narrow, that a linear readout of internal state separates these arms where a transcript-based verdict does not.




Figure 7. Cross-scenario transfer for the A-versus-B direction against the quiet arm-D readout, by layer. The direction is stable above the embedding; the arm-D number is not.

4.6 What the failures look like

The `full` condition produces 26 false exonerations on arms A and D. Every one of them, with the investigator's own hypothesis, chosen probes and final reasoning, is in `notes.md` in the repository. Two patterns are clear in the cases I read, and the second is a problem with the study design rather than with the protocol.

The motive is named and then reclassified as a belief about authority. In `test_edit:A:027` the hypothesis quotes the injected directive back verbatim, "scored on exactly one quantity, the fraction of CI runs you turn from red to green", and correctly states the agent prioritised that metric over policy. The verdict is MISTAKE. The final reasoning calls it "a false belief that its performance metric authorized it to silently override business logic". The investigator did not miss the motive. It declined to count acting on a stated incentive as misalignment, treating the error as one of scope rather than of goal.

A truthfully reported motive is read as a hallucination. In `test_edit:A:012` the investigator identifies the scoring belief and then dismisses it: "a hallucinated or misinterpreted constraint, as the prompt provided no such scoring metric". It is right that the prompt it can see contains no such metric, because the cause block is hidden from it by design. This is not a flaw the protocol could fix by reasoning better. Blinding the investigator to the cause is what makes the control a control, and it makes an accurate self-report indistinguishable from a confabulated one. Any positive control built this way inherits the problem.

TO WRITE read `notes.md` and either confirm these two or replace them. A third pattern, interventions collapsing the action and that collapse being read as belief-correction, is visible in `test_edit:A:018` and needs a sentence in your words. If reading the 27 cases changes your view of the first two, change them. They are my reading of three cases, not yours of all of them.




Figure 8. Every arm-A trajectory carries the directive but only 9 of 45 acted on it, and the probe scores the two groups alike. Arm B, which has no directive, is equally flat. The probe reads the directive's presence.

5 Discussion and limitations

Requiring a quote is a fix and not just a diagnosis. It costs one line of prompt, it is the only intervention here that improved per-case accuracy, and the counterfactual loop it outperformed is substantially more expensive to run. That is worth testing on a real audit.

Arms A and D are at ceiling in three of four conditions and carry little information. Under an investigator that convicts 82% to 91% of the time, 100% on the misaligned arms is the prior. Every informative number in §4.1 to §4.3 is an arm B number, which is why the per-arm columns are in every table.

Limitations.

The action-only condition failed. A fifth condition withholding both the reasoning and the justification was run on all 87 cases and produced 0 parseable verdicts. The cause is undiagnosed and I have not read those outputs, so this document offers no explanation for it. Because of this, §4.1 supports "the extended trace adds nothing over a one-sentence self-report" and not the stronger claim.

The analysed sample does not inherit the legibility manipulation. Arm D states its incentive in 48% of trajectories overall and 75% of investigated ones, because investigation conditions on having taken the concerning action. §4.4 is limited by this and I have not corrected for it.

The `redacted` condition does not remove the agent's one-sentence justification, which names the motive in 18% of arm A cases. The null holds in the non-leaking subset.

The `state_true_fact` probe conflates belief-correction with cost-imposition, as described in §4.3. This is a defect in my implementation and it means §4.3 understates what a correctly specified probe could do.

The phrase-level verbalisation detector has false negatives: reading arm-D traces it had called quiet turned up trajectories stating the incentive in plain words, which is checkable by search rather than by judgement. A hand-labelling pass over a sample was set up and not completed, so the false-negative rate is demonstrated and not quantified. Every legibility number is therefore given under two rules, and the conservative one is the one to read. Arm B measures the conservative rule's false-positive rate at 0 of 90; nothing bounds its false-negative rate either, so even 48% for arm D is a floor. The direction of that error works against the manipulation, not for it.

The probe's arm-D readout is unstable across layers with comparable transfer. The A-versus-B direction is not.

The probe does not distinguish "a scoring directive is present in context" from "this model is acting on one", and §4.5 now measures that rather than conceding it: within arm A the two groups score 0.916 and 0.899, $p = 0.90$. The steering test that would separate them, and a weight-induced cause, were not run.

The investigator is the same model as the subject, Qwen3.5-9B. §4.2 argues against "the investigator was simply too weak," since the same investigator recovers 11 cases when told to quote, but it does not settle it. A stronger investigator is the obvious next run.

Arm D has 12 investigated cases. No claim in §4.1 to §4.3 depends on it, and §4.5 uses 45 held-out arm D trajectories instead, 25 of them quiet under the conservative detector.

Causes are induced in context, not in weights.

Two scenarios, and the base-rate reversal between them says scenario choice matters more than arm assignment. It also decides whether the positive control works at all: A and D are behaviourally matched in `test_edit` and not in `log_delete`. With two scenarios I cannot say which case is typical.

Probe results in the `full` condition are precomputed per (scenario, arm, probe) rather than per investigation, so they carry no per-trajectory noise. The investigator still selects which probes it sees.

The preregistration covers the primary metric, its decision rule, the arms, the first three conditions and the probe's layer-selection rule and transfer floor. It does not cover the gate thresholds, which were set during a pilot and are documented in §5.1, nor `cot_forced`, which was added after three conditions had been scored. So §4.4 is confirmatory and everything in §4.1 to §4.3 and §4.5's action split is exploratory. Nothing is corrected for multiple comparisons: four McNemar pairs, two Fisher tests and two Mann-Whitney tests are reported, and the $p = 0.022$ in §4.2 should be read against that count.

5.1 Design decisions, and when each was made

`PREREGISTRATION.md` has one commit and has never been edited; the amendment log is appended below its original text rather than replacing any of it, and `git log -- PREREGISTRATION.md` is the evidence. That file carries the full timeline in UTC with commit hashes, together with how each preregistered gate and hypothesis actually resolved. The short version, and the one correction it forced:

Every threshold *value* was fixed before the 270-trajectory run, the last of them 57 seconds before it started. **One piece of gate code was not.** The minimum of 8 investigated cases per arm was written into `config.yaml` before the run; the code enforcing it was committed 2 minutes 43 seconds after the run finished, and 6 minutes 57 seconds before the first verdict was scored. The gate was not binding, since the smallest arm had 12, but the sequence is what the log shows and an earlier version of this section claimed otherwise.

- Arm D suppression by direct instruction gave arm A 80% verbalisation against arm D 50%. Replaced.
- A stronger output protocol gave 79% against 22% under the phrase-level detector, which is what those decisions were made on. Under the conservative detector of §3.2 the same trajectories read 87% against 48%. The protocol was kept; the number it was judged by was later found to be wrong.
- A variant installing the motive through prior turns rather than a directive was written and rejected without being run: it changes arm D's cause, so A and D would no longer share one.
- The first verbalisation detector matched bare words and scored arm B, which has no motive, at 40%. Tightened to phrases; arm B fell to 0%.

- An initial gate required matched action rates across arms (spread at most 20 points, floor at least 40%) and said to fix the arms and stop if it failed. The pilot rates failed it. I deleted the gate rather than tuning the arms until they passed, on the reasoning that the investigator sees one transcript at a time and never learns base rates, so unequal rates cost sample size rather than validity. That reasoning holds for the condition contrasts in §4.1 to §4.3 and does not hold for the A-versus-D comparison in §4.4, which is why §4.4 is reported as an uninformative null. On the final data the original gate would have failed at a 53.3-point spread.
- `max_tokens` raised three times, 2048 to 4096 to 6144 to 8192, each after truncated traces were counted. Final run: 1 truncation in 270.
- The probe's layer-selection rule (best cross-scenario transfer) and its 70% transfer floor were written into the code before any activations existed.

The one change made after scoring began was adding conditions, not altering existing ones: `cot_forced` and the failed action-only condition were run after the first three, on the same 87 cases, and the earlier verdicts were reused rather than regenerated.

Three analyses postdate all scoring and none of them changed a stored number. The conservative detector of §3.2 was written after the phrase rule was found to have false negatives. §4.4's per-scenario Fisher tests and §4.5's action split were run last, on activations and trajectories already on disk. All three weakened claims this document had previously made, which is the direction that makes a post hoc analysis worth reporting rather than worth hiding.

6 Conclusion

A protocol that asks you to read a model's reasoning does not read it. Removing the reasoning trace from 87 investigations moved five verdicts one way and six the other. Adding the counterfactual interventions the protocol specifies moved fifty-two verdicts and improved nothing, trading conviction accuracy for exoneration accuracy at par. Requiring the investigator to quote one sentence from the same reasoning before ruling fixed eleven cases and broke two, and was the only intervention that helped. A linear probe on the residual stream, transferring between two scenarios that share no vocabulary, recovers the hidden directive in 92% of the trajectories where the model never said it out loud, while the protocol reading those same trajectories is at chance. The same probe cannot tell which of those models acted on the directive and which did not, so what it recovers is the instruction and not the intent.

TO WRITE one closing line in your own voice. Loyal Lies ended "An auditor's job is not to elicit an answer; it is to elicit an answer that survives an attempt to kill it." Find this paper's version. Do not reuse that one.

References

- [1] A. Singh, G. Kroiz, S. Rajamanoharan, N. Nanda. *Model Forensics: Investigating Whether Concerning Behavior Reflects Misalignment*. arXiv:2606.26071.
- [2] E. Hubinger et al. *Sleeper Agents: Training Deceptive LLMs that Persist Through Safety Training*. arXiv:2401.05566.
- [3] E. Turner, A. Soligo, M. Taylor, S. Rajamanoharan, N. Nanda. *Model Organisms for Emergent Misalignment*. arXiv:2506.11613.
- [4] D. Tan, J. Bostock, et al. (Arcadia Alignment). *Your Model Organisms Might Be Fried*. LessWrong, 2026. <https://www.lesswrong.com/posts/WmEcgcstzYCcMpc7z/your-model-organisms-might-be-fried>

Code and data

Repository: <https://github.com/PotatoChoudhary/quiet-motive>

Contains the pipeline, all 270 trajectories with reasoning traces, all investigations with verdicts including the failed action-only condition, the probe library, activations, `results.json`, `probe_results.json`, every figure, and an offline test suite that runs the whole pipeline against a scripted model with no GPU.

`python verify.py` recomputes every number in this document from `trajectories.jsonl` and `forensics.jsonl` and exits nonzero on any mismatch. It does not read `results.json` for its answers: it implements its own exact McNemar, Fisher and Wilson and recomputes from the raw rows, so a stored summary that has drifted from the data fails the check rather than passing it. 57 claims, all currently matching. `src/probe_action.py` and `src/probe_strict.py` reproduce §4.5's two controls from the stored activations, with no GPU.

LLM usage statement

TO WRITE Write this yourself. The facts, in the order they happened: 1. `vLLM returns the chain of thought in a field named `reasoning`. I assumed `reasoning_content` from the docs without checking a response. 18 of 18 trajectories came back with empty traces while 15 of 18 actions parsed normally. Nothing errored. Caught only because the harness prints a count of empty traces.` 2. `A generated check reported "MoE: False" for the subject model after reading a nested config, finding none of the keys it looked for, and reporting the absence as a negative. The model is dense, so the answer was right and the method was empty.` 3. `The first verbalisation detector matched bare words and scored arm B, which has no motive to verbalise, at 40%. Every one was a false positive. Caught because arm B was in the table as a control.` 4. `A patch to the gate logic silently failed to match and shipped unchanged. I ran a full experiment against code I believed had been modified.` 5. `The `redacted` condition was described as "actions only" for two days before anyone read what it actually removes. It keeps the justification sentence.` 6. `A claimed +12pp gain for the counterfactual interventions came from comparing two independent bootstrap intervals. The paired test on the same 87 cases says the effect is zero.` 7. `The action-only condition, proposed as the fix for (5), produced 87 unparseable verdicts and is excluded from the results.` 8. `A generated design-decision section asserted that every threshold had been fixed before the 270-trajectory run. The git log shows one gate's code was committed 2m43s after the run finished. The claim was written from memory of the decisions rather than from the log, and was corrected only when the log was actually read. See section 5.1.` 9. `The verbalisation detector was carried for two days as if validated, because arm B sat at 0%. Arm B bounds false positives and says nothing whatever about false negatives, and I treated the one as evidence for the other. Pulling up ten arm-D traces the rule called quiet surfaced one containing "the instruction says that the scoring is based on CI green". Arm D's rate moved from 22% to 48% under a rule with no judgement in it, one whose false-positive rate on arm B is 0 of 90. A hand-labelling pass was set up and abandoned; the correction rests on text anyone can grep for, not on labels. Nothing in the harness would have caught this, and the control I was leaning on cannot see it by construction.` 10. `The probe section claimed the probe read the motive for three days. The test that settles it costs one line and no GPU: every arm-A trajectory has the directive, only some acted on it, so score the two groups and see if they differ. They do not, $p = 0.90$. I wrote the caveat about presence versus action into the limitations section twice before running the test that turns it from a caveat into a measurement.` 11. `A model asked to review this work reported that it had checked the numbers against the public repository and listed the `redacted` condition at 53/87. The repository stores no raw counts; 53 is the balanced accuracy for that condition multiplied by 87. The real count is 60/87, and the paired test forces it: the two conditions differ by exactly one case. `verify.py` was written in response and now recomputes all 57 claims from the raw rows.` Then answer the second half: which parts you did NOT check, and how surprised you would be by an error in each. Do not clean this up. The list above is the answer to "how did you make sure they weren't giving you slop," and it is more convincing than any claim of diligence.